

F2FS and EXT4 Reliability

Michael Meerbott

Abstract

Flash storage is becoming more prevalent in user level devices. Today, the Android operating system defaults to using EXT4 as its default filesystem, as do many Linux distributions. Most mobile phones use NAND flash storage devices and many laptops today are using the same. F2FS was created by Samsung to take advantage of the NAND flash structure to optimize speeds [1]. In this project, I wanted to compare the reliability of F2FS to that of EXT4 to help people, including myself, decide if it is worth the risk to switch their flash storage systems over to F2FS. I age both EXT4 and F2FS and show that F2FS does not seem to outperform EXT4.

1. INTRODUCTION

Over time, filesystems may operate slow due to file fragmentation. Filesystem aging is the measure of fragmentation over time. In this project, I use the aging method from the paper “File Systems Fated for Senescence? Nonsense, Says Science!” by A. Conway et al. [2]. The method uses a sequence of Git pulls of the Linux source code onto a device. After a number of pulls, it then performs a grep test. The grep test simply performs a grep of a non-existing string on the root directory of the source code and times how long the grep takes. It then continues on, creating a timeline of grep times after a total number of pulls.

2. Methodology

System Setup

The project was performed on a ThinkPad E330 running Ubuntu 16.04.5 (Linux 4.4.0-104-generic). The tools for F2FS used, F2FS-tools mkfs-f2fs, was 1.6.1. The filesystems were aged on a 32 GB SanDisk Ultra class 10 SD card that was split evenly into 4 partitions of 7.43 GB each. The card used 2 partitions of F2FS and 2 partitions of EXT4. Both filesystems use one partition for aging and one “clean” partition to compare the aging against. Below are the configurations I used to create the filesystems, which are the defaults:

F2FS

```
batched_trim_sections:32
cp_interval:60
dir_level:0
gc_idle:0
gc_max_sleep_time:60000
gc_min_sleep_time:30000
gc_no_gc_sleep_time:300000
ipu_policy:16
```

```
max_small_discards:0
max_victim_search:4096
min_fsync_blocks:8
min_ipu_util:70
ram_thresh:10
ra_nid_pages:4
reclaim_segments:755
```

EXT4

```
rw
block_validity
delalloc
barrier
user_xattr
acl
resuid=0
resgid=0
errors=continue
commit=5
min_batch_time=0
max_batch_time=15000
stripe=0
data=ordered
inode_readahead_blks=32
init_itable=10
max_dir_size_kb=0
```

Git Pull Configurations

There are 2 configurations used for the git pull aging. The first tests run 10,000 pulls and perform grep tests every 100 pulls. The second set runs 5,000 pulls and perform grep tests every 50 pulls.

3. Issues and Difficulties

The git-pull script was simple to set up, but didn't run as smoothly as hoped. After around 200 git pulls, the script would crash on a git pull with error 128. I re-ran the script several times with the first set of parameters. I looked around online for error 128, but the search results were not consistent. Many of the search results were project specific (a project unrelated to what I was doing), and the others were mentioning a bad git ssh key, which was not my issue. The script would stop running after the first error, so I decided to simply insert a try-except block to see if the script would continue immediately after the first error. I let it run for a few days, and when I checked the results, the script had not continued after the error immediately, but instead after many more errors, and would fail again. It was constantly off and on. I ran it 3 times with the 10,000-pulls setting to make sure it would keep running with occasional results, and it seemed they were all at random pulls.

I decided 10,000 runs was going to take too long with too little results, so I reduced the tests to 5,000 pulls and

grep tests after every 50 pulls. I ran the script again for a few days for both filesystems, but the results did not show any notable aging, so I ran them again: 1 week for the F2FS tests and 1 week for the EXT4 tests. Those are the results use in this report.

4. Results

These results do not have a constant step size between the grep tests. This is due to the failed git pulls, so I only added the pulls that succeeded. In Figure 1, we see that F2FS shows some aging after 250 git pulls. At the last grep test (393 pulls), the results show the aged partition is about 1.14x slower than the unaged partition. Note that the line should be compared with the non-aged line as a jump in Grep Cost can mean an increase in number of files. The results also show some "free play" range: in the initial pull, both partitions are clean and unaged, yet we see a difference of .677 sec/GB in the grep cost.

Figure 2 shows our results for EXT4. These results do not show any aging, since the difference between all of the points is close to the "free play" we see on the initial test. The difference in the initial test is .481 sec/GB. These results tell us that EXT4 is more reliable in terms of aging than F2FS since it showed less aging.

Figure 3 compares the F2FS and EXT4 results against each other. Surprisingly, the initial EXT4 results outperform F2FS by around 50% (EXT4 is 1.5x faster). SD cards use NAND flash storage, so F2FS was expected to run faster. I am not sure why EXT4 outperformed F2FS; it could be the version of f2fs-tools I used, or the version of Linux I used.

Overall, I would say that these results are inconclusive. There is not enough data/git-pulls to decide anything. For conclusive results, I would expect to see some obvious EXT4 aging. The fact that EXT4 immediately outperformed F2FS is also odd, especially since most speed tests show F2FS performing better.

5. Future Work

I would recommend using more than one aging tool for this and testing it on different storage devices or systems since my laptop had EXT4 always outperforming F2FS. A better method should also be considered, since the script used goes by commits in one project. This would have a limited number of changes per commit pulled. This may not provide enough data changes for aging the filesystem.

References

- [1] Changman Lee, Dongho Sim, Joo-Young Hwang, and Sangyeun Cho. 2015. F2FS: a new file system for flash storage. In Proceedings of the 13th USENIX Conference on File and Storage Technologies (FAST'15). USENIX Association, Berkeley, CA, USA, 273-286.
- [2] Alex Conway, Ainesh Bakshi, Yizheng Jiao, Yang Zhan, Michael A. Bender, William Jannen, Rob Johnson, Bradley C. Kuzmaul, Donald E. Porter, Jun Yuan, and Martin Farach-Colton. 2017. File systems fated for senescence? nonsense, says science!. In Proceedings of the 15th Usenix Conference on File and Storage Technologies (FAST'17). USENIX Association, Berkeley, CA, USA, 45-58

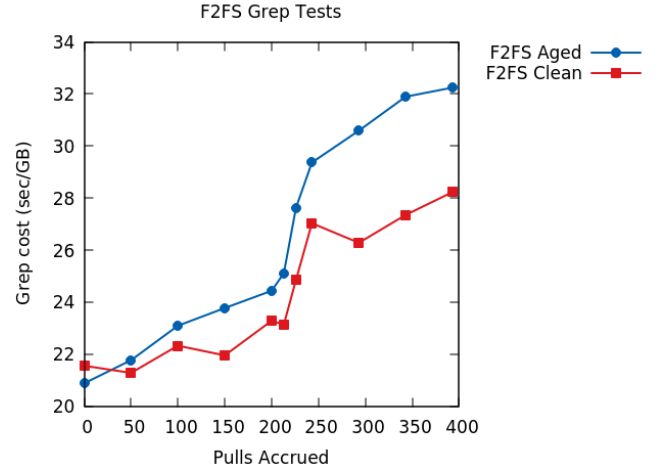


Fig. 1. The F2FS results with aged and unaged tests

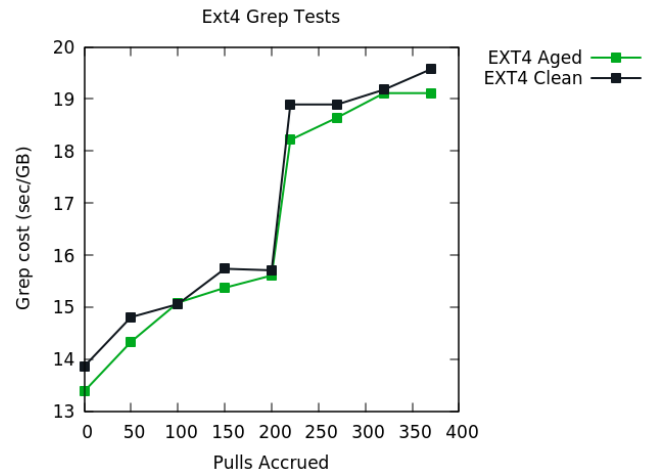


Fig. 2. The EXT4 results with aged and unaged tests




Fig. 3. Both F2FS and EXT4 (aged and unaged) results compared